

Holonomy-Aware Topological AlphaZero: Self-Play Go on Non-Orientable and Arbitrary-Genus Boards

Sam Leventhal

Kahlert School of Computing, University of Utah*

July 8, 2026

Abstract

Game-playing agents assume flat rectangular boards: the convolutional towers of AlphaZero and KataGo are defined over $\mathbb{Z}^2$ tensors and cannot even ingest a board on a sphere, a torus, or a Möbius band. We introduce *Geodesics*, a topological generalization of Go to arbitrary closed and bordered surfaces (spheres, tori, cylinders, Möbius bands, Klein bottles, real projective planes) and lattices in one and three dimensions, under any of several mesh families — and ask what architecture plays it optimally. We present a first implementation of a *Holonomy-Aware Topological AlphaZero* (HATZ): an AlphaZero-style self-play agent whose network (i) carries the reflection ($\mathbb{Z}_2$) part of the gauge structure that non-orientable surfaces provably require, via an orientation cocycle computed from the board’s combinatorics and sheaf-style transport maps conditioned on it; (ii) is exactly equivariant to each board’s automorphism group by construction, replacing the dihedral data augmentation that has no analogue off the square grid; (iii) pools through a Morse–Smale multi-persistence hierarchy of the position’s influence field, a canonical topology-grounded coarsening; and (iv) reads the board as a cell complex, passing messages over vertex–edge–face incidences with curvature and localized-homology inputs and an ownership auxiliary head. The network is small (3.5×10^4 parameters), dependency-free, and every gradient is hand-derived and numerically verified. We verify the orientation cocycle against ground-truth orientability on eleven board types, demonstrate machine-precision equivariance, and report two self-play pathologies specific to this setting (transposition-table ko cycles; a degenerate early-pass equilibrium under Tromp–Taylor scoring) together with their remedies. To our knowledge this is the first game-playing agent on non-orientable boards; a full transfer and ablation study is specified and in progress.

1 Introduction

The rules of Go never needed a flat grid. Stones occupy vertices of a graph; liberties are empty neighbors; a chain lives while its liberty set is non-empty; territory is graph reachability. Everything downstream of these definitions — capture, ko, life and death, Tromp–Taylor scoring — is intrinsically combinatorial. The 19×19 lattice is a historical accident of the plane. *Geodesics* takes the rules at their word and instantiates them on discrete surfaces: geodesic and Goldberg

*Draft v0.1. Code, boards, and the browser-playable implementation: <https://samleventhal.com/subpages/geodesics/geodesics.html>.

spheres, square and hexagonal tori, Möbius bands, Klein bottles, the real projective plane, cylinders, and cubic lattices in E^3 , with vertex, edge, face, or full cell-complex incidence structures.

This generality breaks every assumption baked into modern game networks. AlphaZero [1] and KataGo [2] use residual convolutional towers over rectangular tensors; their 8-fold dihedral data augmentation exploits a symmetry group that simply does not exist on a sphere (whose geodesic boards have icosahedral symmetry) or a Klein bottle (whose isometries include orientation-reversing translations). More fundamentally, Weiler et al. [13] prove that convolution on a non-orientable surface cannot reduce its structure group to $SO(2)$: any consistent feature transport must be steerable under reflections, because parallel transport around an orientation-reversing loop returns frames mirrored. A network that ignores this is discontinuous across the orientation seam of a Möbius or Klein board.

We ask: *what is the right architecture for optimal play across domain geometries and topologies*, and we build a first honest draft of the answer. Our contributions:

1. **The first game-playing agent on non-orientable and arbitrary-genus boards.** We compute a representative orientation cocycle $\varepsilon \in Z^1(G; \mathbb{Z}_2)$ of the first Stiefel–Whitney class from the board’s combinatorics alone (dual-BFS over face cycles when faces exist; the quotient-gluing rule for fundamental-domain grids), verify it against ground truth on eleven board types, and pass messages through ε -conditioned transport maps so that crossing the seam applies a genuinely different learned map (§4.1).
2. **Exact automorphism equivariance without augmentation.** Weight-shared message passing over the board graph commutes with every graph automorphism; we keep all inputs Aut-invariant by construction and verify equivariance to machine precision, including under the seam-preserving isometries of the Möbius band (§4.5). This replaces dataset-specific augmentation with intrinsic symmetry, with the strict generalization benefit of equivariance [16].
3. **Morse–Smale multi-persistence pooling.** The ascending and descending basin hierarchy of the position’s diffused influence field [17] segments the board into territories-and-frontiers; we use it as a position-dependent, parameter-independent pooling operator inside the network — a canonical, topology-grounded coarsening where learned pooling such as DiffPool [15] only approximates one, and one through which backpropagation remains exact (§4.4).
4. **A cell-complex, ownership-supervised value/policy net** with curvature (angle-defect) and localized-homology inputs, KataGo-style global pooling, and an ownership head; 3.5×10^4 parameters, pure-`numpy` with hand-derived gradients verified numerically to 5×10^{-6} over every parameter (§4.3).
5. **Two self-play pathologies and their fixes:** PUCT descends that cycle forever through ko transpositions when the in-tree engine forgets history, cured by a positional-superko guard; and a degenerate early-pass equilibrium under Tromp–Taylor-plus-komi scoring that collapses the policy, cured by masking passes early in self-play (§5).

The literature intersection we occupy — game playing $\times$ gauge equivariance $\times$ non-orientable surfaces — is, to our knowledge, empty.

2 Related work

Game networks and size generality. KataGo trains one convolutional net across board sizes with masking and global-pooling bias layers and extrapolates beyond its training sizes [2]; Polygames obtains board-size invariance from fully convolutional trunks with global pooling [3]; ScalableAlphaZero replaces the CNN with a GIN and transfers from small to large Othello boards with an order-of-magnitude training-time reduction [4]. All operate on rectangular (or at least fixed-family) boards; none addresses topology.

GNN expressivity. Message passing is bounded by the 1-WL test [5, 6]; Go meshes are locally regular, which is precisely where 1-WL is weakest, motivating multi-aggregator updates [7], higher-order structure, and our cell-complex lifting, which is strictly more expressive than 1-WL [8].

Gauge equivariance. Gauge-equivariant CNNs on manifolds and meshes [11, 12] achieve equivariance to local frame choice via parallel transport; Weiler et al. [13] prove non-orientable surfaces force the structure group $O(2)$ (reflections included), demonstrating on a Möbius toy model. We implement the discrete reflection ($\mathbb{Z}_2$) part, which is the component non-orientability makes mandatory.

Topological deep learning. Simplicial and CW networks [8], combinatorial complex networks [9], and neural sheaf diffusion [10] pass messages over incidence structures with (in the sheaf case) learnable transport. Our boards *are* cell complexes, and our transport maps are a transfer-friendly sheaf: globally shared matrices conditioned on the edge’s cocycle value, rather than per-edge free parameters.

Hierarchical pooling and topological segmentation. Graph U-Nets [14] and Diff-Pool [15] learn coarsenings; our pooling is the Morse–Smale multi-persistence hierarchy of a scalar field on the board, following the hierarchical GNN construction of [17], which gives a canonical segmentation into basins (territories) meeting along separatrices (contested frontiers).

3 Setting: Go on a surface

A *board* is a finite graph $G = (V, E)$, optionally equipped with a 2-cell structure (face cycles) making it a closed or bordered combinatorial surface, produced by `make_board(surface, mesh, resolution)`. The surfaces used here, with the smallest instances in our suite:

board	$ V $	$ E $	χ	orientable	seam edges
sphere (geodesic $f=2$)	42	120	2	yes	0
plane D^2 (5×5)	25	40	1	yes	0
cylinder (5×5)	25	45	0	yes	0
torus (5×5)	25	50	0	yes	0
Möbius (5×5)	25	45	0	no	5
Klein (5×5)	25	50	0	no	5
$\mathbb{RP}^2$ (5×5)	25	48	1	no	8
box E^3 (4^3)	64	144	—	—	0
3-torus (4^3)	64	192	—	—	0

Rules are graph-intrinsic (positional superko; Tromp–Taylor scoring with komi 7.5). Quotient boards are fundamental-domain grids with identifications; e.g. the Klein bottle identifies $(x+n_x, y) \sim (x, n_y-1-y)$ (a flip) and $(x, y+n_y) \sim (x, y)$; $\mathbb{RP}^2$ flips both pairs. The browser

implementation and the training package construct adjacency independently and agree edge-for-edge (verified in tests), so trained agents transfer between them.

4 Architecture

4.1 Reflection gauge structure from combinatorics

A local orientation at a vertex is a $\mathbb{Z}_2$ gauge choice; an assignment $\varepsilon : E \rightarrow \{\pm 1\}$ records whether traversing an edge preserves or reverses it. Gauge transformations act by $\varepsilon'(u, v) = g(u) \varepsilon(u, v) g(v)$ for $g : V \rightarrow \{\pm 1\}$; the *holonomy* of a cycle, $\prod_{e \in \gamma} \varepsilon(e)$, is gauge-invariant, and the class $[\varepsilon] \in H^1(G; \mathbb{Z}_2)$ restricted to surface cycles represents the first Stiefel–Whitney class w_1 : it is trivial iff the surface is orientable.

We compute a representative in two ways. When face cycles are available, a BFS over the dual graph assigns each face an orientation; a shared edge whose two incident faces traverse it incoherently under the assignment is a *seam* edge ($\varepsilon = -1$). For fundamental-domain grids the flipped gluing itself is the seam: x -wrap edges on Möbius and Klein boards, both wrap families on $\mathbb{RP}^2$.

Proposition 1. *The seam set produced by dual-BFS is a cocycle representing w_1 ; it is empty for some BFS iff the surface is orientable. On quotient grids, the flipped-gluing edge set is likewise a representative.*

Proof sketch. Dual-BFS fixes a gauge on faces along a spanning tree; conflicts can occur only on non-tree dual edges, and the conflict indicator is exactly the obstruction cocycle to extending the orientation, whose class is w_1 . For quotient grids, cutting along the flipped gluing yields the orientable fundamental domain, so a loop reverses orientation iff it crosses the gluing an odd number of times. $\square$

Empirically, the construction recovers ground-truth orientability on all eleven board types in our suite (including hexagonal-mesh Möbius bands and 3D lattices), with seam counts matching the gluing combinatorics (Table above).

4.2 ε -conditioned sheaf transport

Let $h_v \in \mathbb{R}^D$ be vertex features. Messages pass through restriction maps conditioned on the cocycle,

$$m_v = \frac{1}{|N(v)|} \sum_{u \in N(v)} (W_0 + \varepsilon(u, v) W_1) h_u = W_0 (Ah)_v + W_1 (A_\varepsilon h)_v, \quad (1)$$

where A is the row-normalized adjacency and A_ε its ε -signed counterpart. Crossing the orientation seam therefore applies a different learned map — the mechanism by which neural sheaf diffusion [10] encodes orientation reversal and heterophily — but with W_0, W_1 *shared globally* rather than learned per edge, so the same parameters transfer across boards, and because ε enters linearly the layer reduces to two fixed aggregation matrices and remains trivially differentiable by hand.

Honest scope. This is the discrete reflection part of the $O(2)$ structure non-orientability mandates [13]; the continuous $SO(2)$ part (anisotropic kernels with rotational transport over the mesh’s rotation system) is future work, and our aggregation is isotropic. Further, the network consumes ε in a fixed canonical gauge; holonomy-derived inputs are gauge-invariant, but the

learned transport is gauge-covariant only up to reparameterization. We therefore train with *gauge augmentation*: each self-play game and each optimization step resamples g uniformly and uses $\varepsilon'(u, v) = g(u)\varepsilon(u, v)g(v)$, teaching the invariance the architecture does not yet guarantee.

4.3 Cell ranks and inputs

The board is read as a complex: vertex features always; edge features always (seam flag, endpoint color pattern); face features when face cycles exist. Per layer ℓ ,

$$H_V \leftarrow \text{relu}(H_V W_s + (A H_V) W_0 + (A_\varepsilon H_V) W_1 + (M_{EV} H_E) W_{ev} + \bar{h} W_g + b), \quad (2)$$

$$H_E \leftarrow \text{relu}(H_E W_{es} + (M_{VE} H_V) W_{ve} + (M_{FE} H_F) W_{fe} + b_e), \quad H_F \leftarrow \text{relu}(\dots), \quad (3)$$

with $M_\cdot$ fixed incidence-mean operators and $\bar{h}$ the global mean (KataGo-style pooling bias [2]). Inputs are Aut-invariant by construction: stone/empty indicators from the mover’s perspective, normalized degree, discrete curvature (the Gauss–Bonnet angle defect $2\pi - \sum_i \theta_i$, whose total is $2\pi\chi$), localized H_0 data (own/opponent group size and liberty counts — the persistent-homology features in their cheapest useful form), and global scalars $[1[w_1 \neq 0], \chi/|V|]$. Heads: a per-vertex policy logit plus a pooled pass logit; a pooled tanh value; and a per-vertex tanh *ownership* head supervised by final Tromp–Taylor territory, KataGo’s auxiliary target adapted to graphs. Total: 34,916 parameters at width 32×3 layers — ~ 140 KB, comfortably client-side. All gradients are hand-derived in `numpy`; central-difference verification over every parameter gives worst relative error 4.5×10^{-6} (Möbius: seam transport + pooling + ownership) and 5.8×10^{-7} (tetrahedron: full vertex–edge–face coupling).

4.4 Morse–Smale multi-persistence pooling

Let f_0 be the two-step mean-diffused stone field of the position. Its ascending and descending basin partitions — computed by the same union-find persistence hierarchy that drives our earlier hierarchical GNN [17] and, incidentally, the game’s built-in engine — segment the board into influence basins joined at separatrices, a decomposition that is Go-semantic (territories and contested frontiers) and respects the surface’s critical structure. At a persistence quantile we form regions r , the one-hot assignment $S \in \{0, 1\}^{|V| \times R}$, mean-pool P , region adjacency A_r , and insert after the first layer

$$Z = \text{relu}((P H_V) W_p + (A_r P H_V) W_r), \quad H_V \leftarrow H_V + (S Z) W_u. \quad (4)$$

The segmentation depends on the *position*, not on parameters, so S, P, A_r are constants of each forward pass and backpropagation through the block is exact — a guarantee learned pooling lacks.

4.5 Equivariance

Proposition 2. *If every input feature is invariant under $\text{Aut}(G)$ (and, where used, under automorphisms of the complex and of ε), the network is exactly Aut-equivariant: for any automorphism σ , $\pi(\sigma \cdot s) = \sigma \cdot \pi(s)$ and $v(\sigma \cdot s) = v(s)$.*

Proof sketch. Each layer is a composition of permutation-equivariant operators ($A, A_\varepsilon, M_\cdot, S, P$ conjugate by the permutation; pointwise nonlinearities and global means commute with it) with shared weights. $\square$

Verified to machine precision: the empty-board policy is constant on the WL orbits of the geodesic sphere (two orbits: the twelve pentavalent and thirty hexavalent vertices; spread 2×10^{-15}) and the torus (vertex-transitive; spread 0), and exactly symmetric under the seam-preserving Möbius reflection (spread 0). On non-orientable boards the *gauge-dependent* remainder — seam placement — is the component gauge augmentation trains away; we report this rather than hide it.

5 Self-play training, and two pathologies

Training is standard AlphaZero: PUCT ($c=1.5$, Dirichlet root noise), visit-count policy targets with temperature 1 for the first ten moves, replay buffer, Adam. Two failures specific to this setting are worth recording.

Transposition ko cycles. Our search merges states in a transposition table keyed by (stones, mover, passes) and re-seats a scratch engine per step, which forgets ko history. Once two ko-mirror nodes exist with PUCT argmaxes pointing at each other, a descent cycles $A \rightarrow B \rightarrow A$ forever — the process hangs, typically from the second iteration once the tree is deep enough to contain a ko. The fix is the search-side analogue of the game’s superko rule: any state repeated within a single descent scores as a draw and terminates (with a depth cap as backstop). Three hundred simulations through a live ko now complete in bounded time; the guard is a permanent regression test.

The early-pass equilibrium. Under Tromp–Taylor scoring with positive komi, an immediate double pass is a *guaranteed White win*. Shallow search discovers this instantly: White’s visits collapse onto pass, half the replay buffer becomes one-hot pass targets, and the learned policy passes as both colors — losing every evaluation game to a random mover. (Loss decreases throughout; the curve looks healthy while the policy degenerates.) KataGo guards its self-play against early passing for the same reason [2]; we mask the pass action from search targets for the first $|V|/2$ moves, after which games run full length and evaluation strength reflects play rather than the scoring artifact.

6 Preliminary verification and evaluation protocol

What is verified now. Cocycle-vs-orientability on 11 board types; full-parameter gradient checks (4.5×10^{-6} , 5.8×10^{-7}); machine-precision equivariance as above; Gauss–Bonnet consistency of the curvature inputs (tetrahedron defect = π per vertex); browser/Python adjacency parity edge-for-edge on cylinder, Klein, and $\mathbb{RP}^2$ boards; end-to-end training and serving (checkpoint $\rightarrow$ WebSocket bridge $\rightarrow$ browser opponent, including seam reconstruction on remote boards). Short training runs on the non-orientable trio (Möbius, Klein, $\mathbb{RP}^2$; 24 simulations, six iterations) show healthy optimization after the pass guard (loss $3.92 \rightarrow 2.81$; parity with a random opponent within noise) — reported as smoke evidence, not results.

Protocol (in progress). (i) *Zero/few-shot transfer*: train on k topologies, evaluate Elo on held-out ones (e.g. train sphere+torus, test Klein), against fixed anchors (random; the hand-designed Morse–Smale engine; a Stage-1 GNN of matched size). (ii) *Ablations*: remove in turn the ε -transport ($W_1=0$), the MSC block, the ownership head, curvature inputs, and gauge augmentation; measure Elo and gauge-sensitivity. (iii) *Ceiling*: on the plane boards only, compare against KataGo as an upper reference. (iv) *Symmetry*: measure sample efficiency versus an augmentation-based baseline, testing the equivariance dividend of [16] in self-play.

7 Limitations and future work

No continuous $SO(2)$ steerability yet (isotropic aggregation over the rotation system); gauge invariance is learned, not guaranteed — a strictly equivariant formulation via the orientation double cover is the natural next step; persistent-homology inputs are localized proxies (H_0 statistics) rather than diagrams, and H_1 features for eyes/ko are future work; automorphism groups are used for verification (WL orbits) but not yet for target symmetrization; and the empirical study above is the actual test of the thesis. A browser export of the network is mechanical — the Morse–Smale machinery already ships in the client-side engine — and would make every trained agent playable with no installation.

8 Conclusion

Go’s rules are indifferent to topology; the field’s architectures are not. We built the first agent that plays where the boards get strange — carrying exactly the reflection structure that non-orientability forces, the automorphism symmetry each board actually has, and a topological pooling hierarchy with guarantees — small enough to ship in a browser tab, honest about which invariances are proved, learned, and pending.

References

- [1] D. Silver et al. A general reinforcement learning algorithm that masters chess, shogi and Go through self-play. *Science* 362(6419), 2018.
- [2] D. J. Wu. Accelerating self-play learning in Go. arXiv:1902.10565, 2019.
- [3] T. Cazenave et al. Polygames: Improved zero learning. arXiv:2001.09832, 2020.
- [4] S. Ben-Assayag and R. El-Yaniv. Train on small, play the large: Scaling up board games with AlphaZero and GNNs. arXiv:2107.08387, 2021.
- [5] K. Xu, W. Hu, J. Leskovec, S. Jegelka. How powerful are graph neural networks? *ICLR*, 2019.
- [6] C. Morris et al. Weisfeiler and Leman go neural: Higher-order graph neural networks. *AAAI*, 2019.
- [7] G. Corso, L. Cavalleri, D. Beaini, P. Liò, P. Veličković. Principal neighbourhood aggregation for graph nets. *NeurIPS*, 2020.
- [8] C. Bodnar et al. Weisfeiler and Lehman go cellular: CW networks. *NeurIPS*, 2021.
- [9] M. Hajij et al. Topological deep learning: Going beyond graph data. arXiv:2206.00606, 2022.
- [10] C. Bodnar, F. Di Giovanni, B. P. Chamberlain, P. Liò, M. M. Bronstein. Neural sheaf diffusion: A topological perspective on heterophily and oversmoothing in GNNs. *NeurIPS*, 2022.
- [11] T. S. Cohen, M. Weiler, B. Kicanaoglu, M. Welling. Gauge equivariant convolutional networks and the icosahedral CNN. *ICML*, 2019.
- [12] P. de Haan, M. Weiler, T. Cohen, M. Welling. Gauge equivariant mesh CNNs. *ICLR*, 2021.
- [13] M. Weiler, P. Forré, E. Verlinde, M. Welling. Coordinate independent convolutional networks — isometry and gauge equivariant convolutions on Riemannian manifolds. arXiv:2106.06020, 2021.

- [14] H. Gao and S. Ji. Graph U-Nets. *ICML*, 2019.
- [15] R. Ying et al. Hierarchical graph representation learning with differentiable pooling. *NeurIPS*, 2018.
- [16] B. Elesedy and S. Zaidi. Provably strict generalisation benefit for equivariant models. *ICML*, 2021.
- [17] S. Leventhal, A. Gyulassy, V. Pascucci, M. Heimann. Exploring classification of topological priors with machine learning: A hierarchical GNN from Morse–Smale segmentation. *NeurIPS Workshop*, 2022.